

# Full Stack Development

## Containers, Microservices and UI

11. Kubernetes




# The Operations Challenge

- In a previous module, we looked at containerization from a development perspective
- In production, we have to be concerned with
  - Coordinating the activity of possibly thousand of containers that need to work together
  - Creating and maintaining connections between containers
  - Ensure the whole system operates well enough to meet Service Level Agreements (SLAs)
- We need to deal with non-functional requirements
  - Loading, throughput, stress, response times
  - Disaster recovery
  - Security
- The lack of an effective way to do this was a major impediment to the deployment of microservice based applications

# Site Reliability Engineering

- Practices designed to ensure large systems are operational
- Continuously checking for potential problems
- Manages a set of mitigation responses to react to problems
- Recent examples
  - Rogers Canada 2022 network failure
  - Facebook October 2021 upgrade failure
  - Check out [risks.org](https://risks.org)
- As applications scale, this becomes increasingly difficult



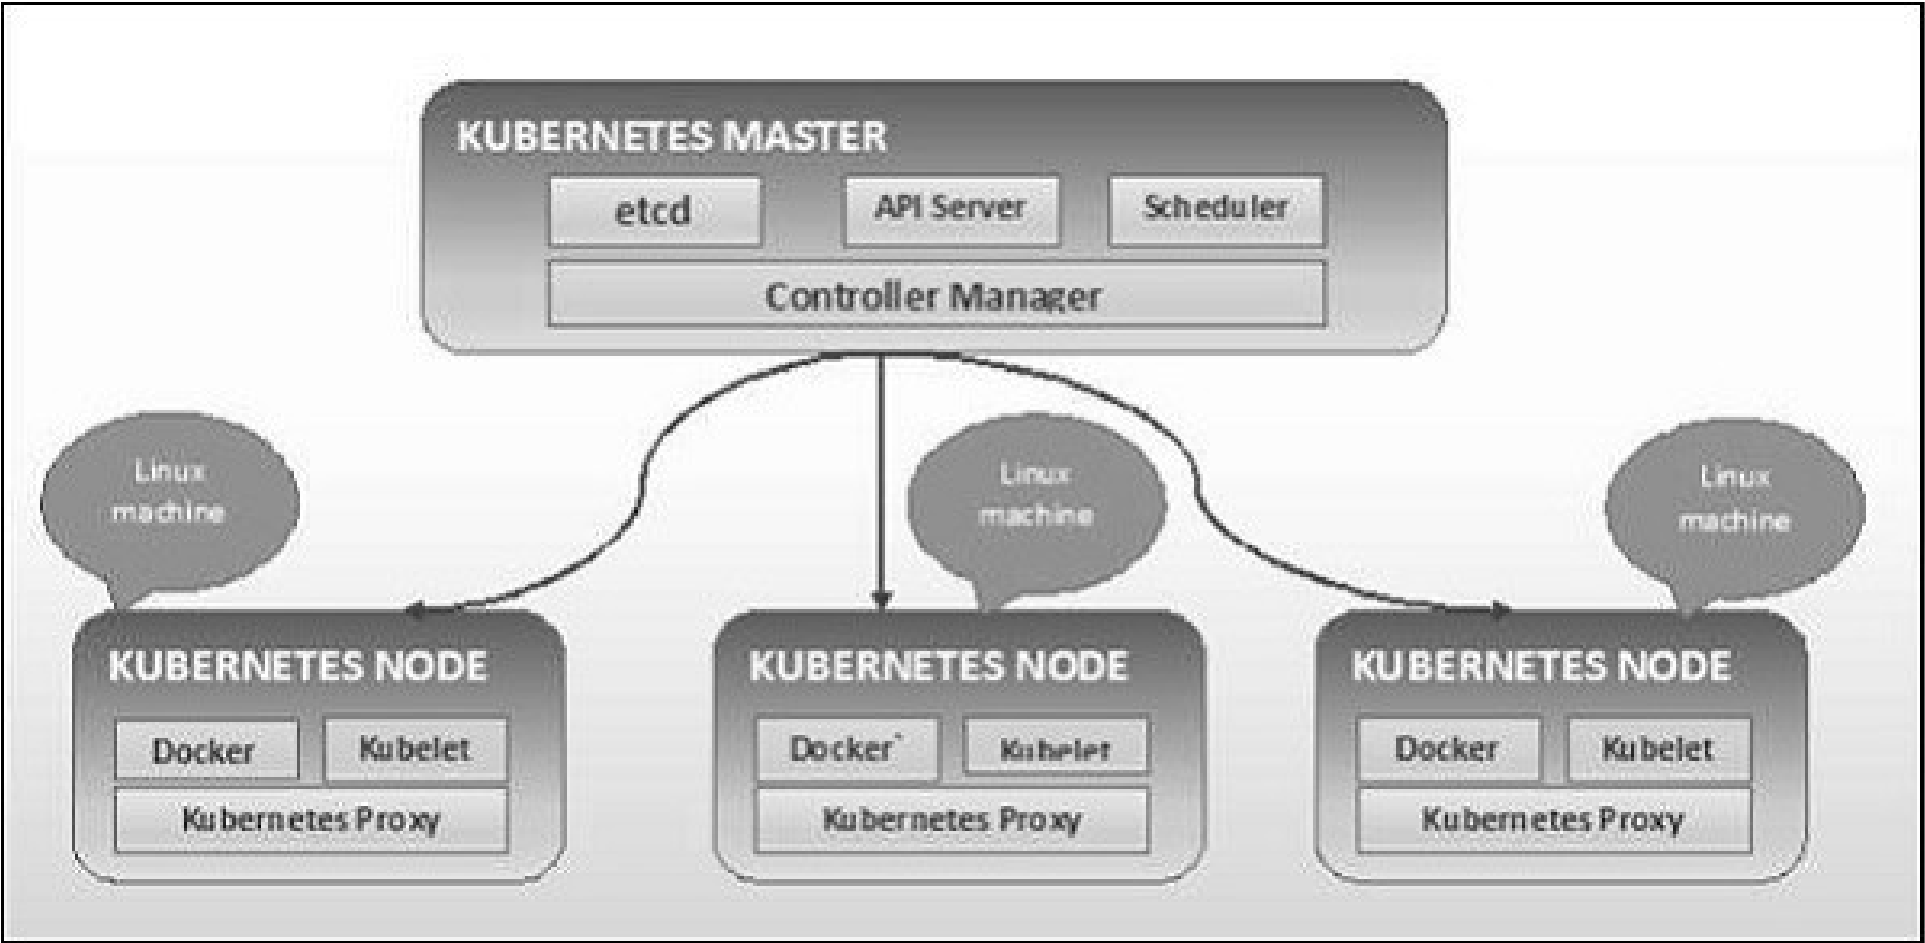
# Kubernetes

- Kubernetes is a container orchestration manager
  - Not the only manager
  - Docker Swarm does the same
  - Kubernetes is “industrial strength”
- Orchestration:
  - Manages “clusters” of containers
  - Provides service discovery
  - Manages scaling and failover
  - Works at the Ops level
  - Infrastructure as Code



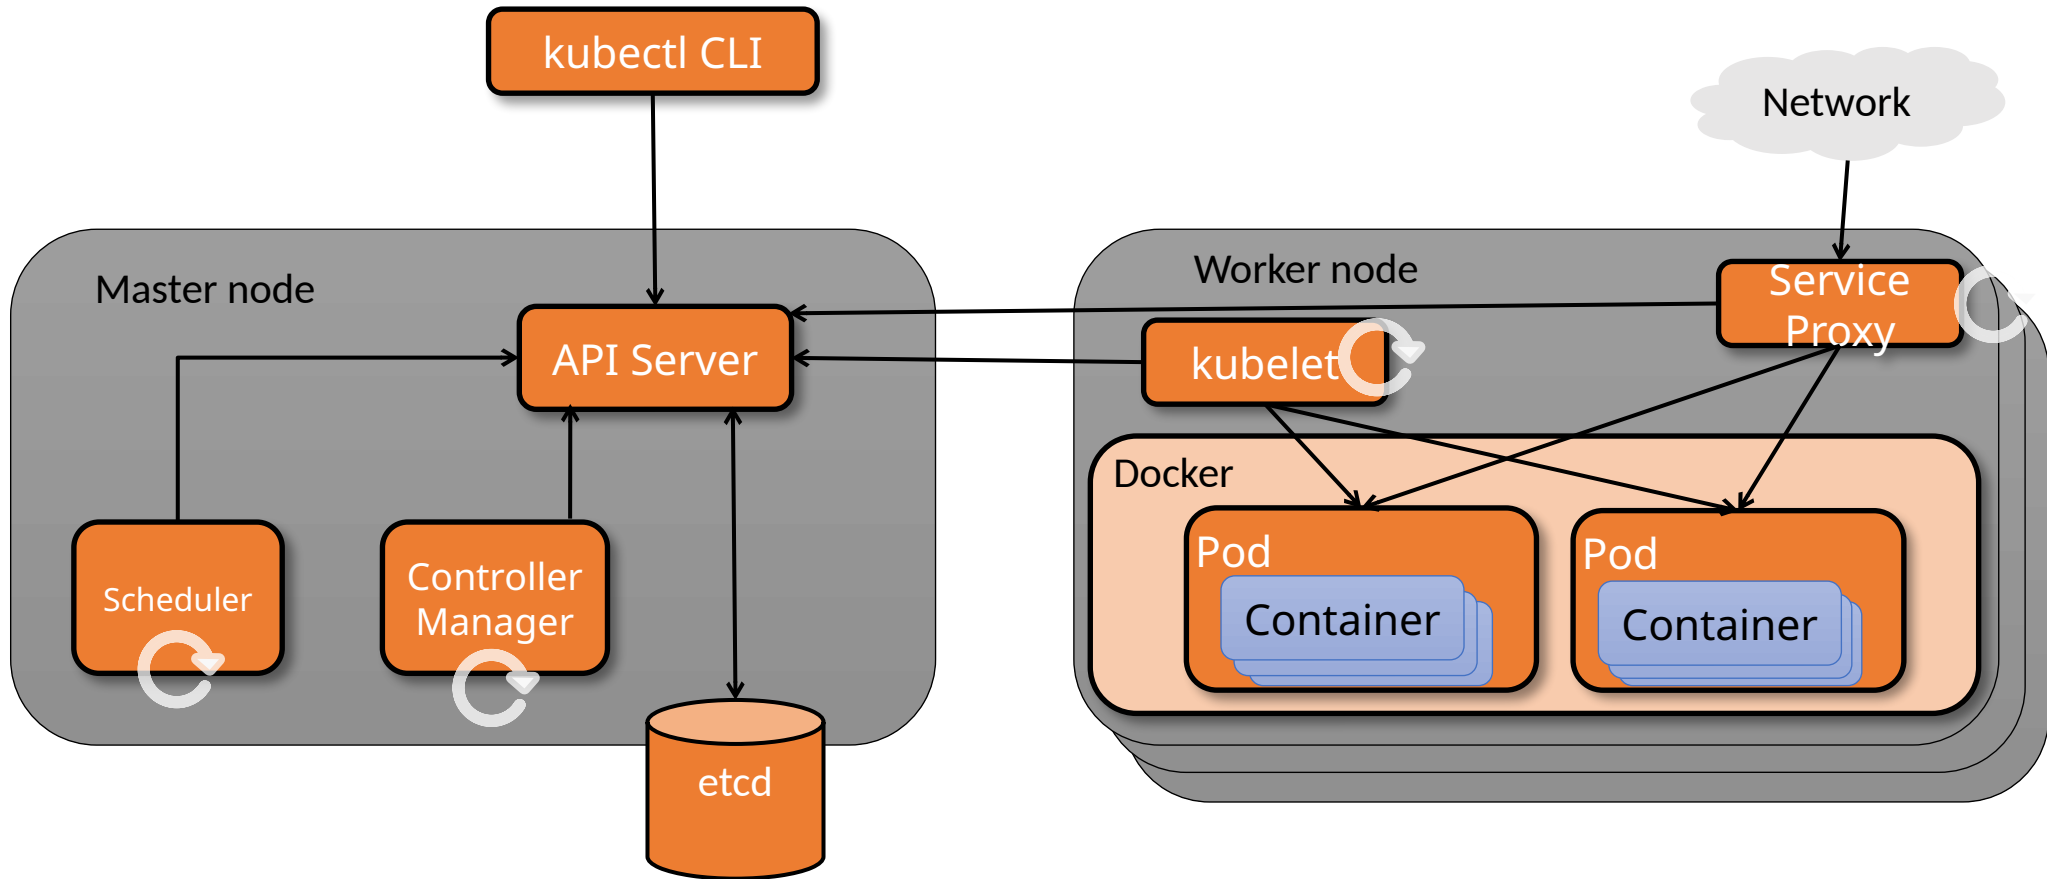
# kubernetes

# Kubernetes Cluster



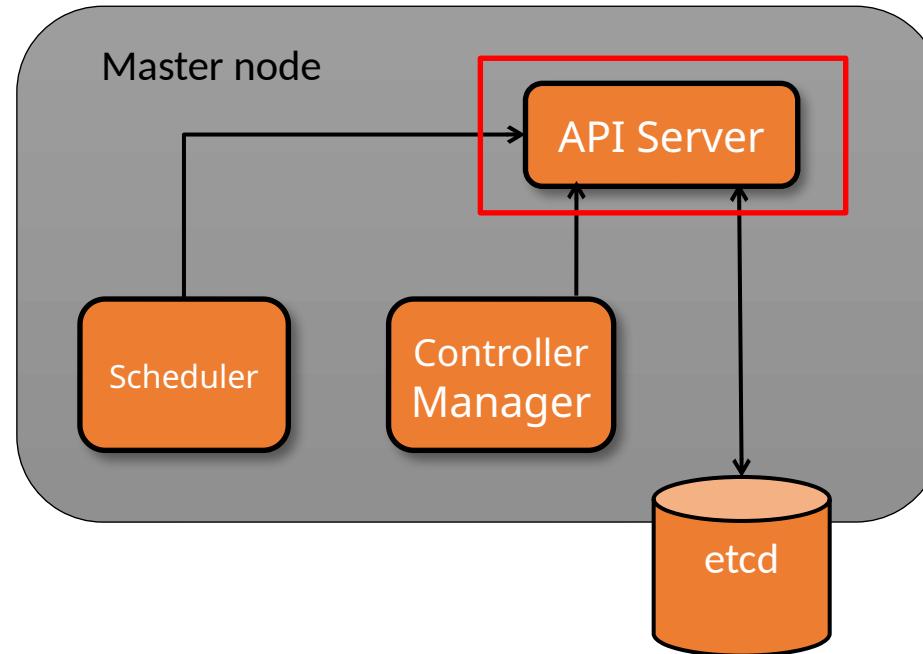
# Kubernetes Architecture

- Kubernetes nodes can be physical hosts or VM's running a container-friendly Linux (kernel > 3.10)



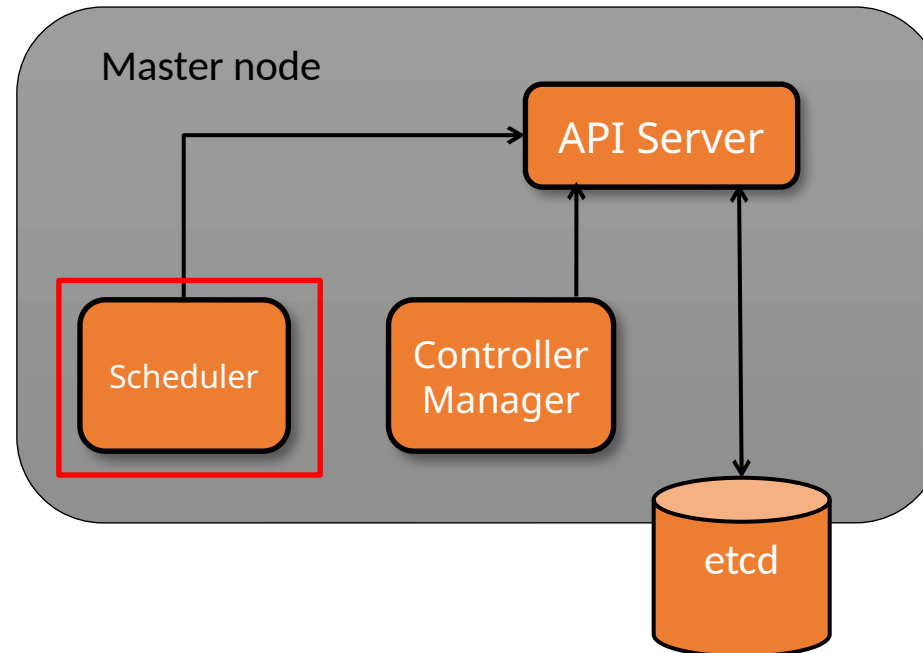
# Master Node Components

- API Server (kube-apiserver): exposes the Kubernetes REST API, and can be scaled horizontally



# Master Node Components

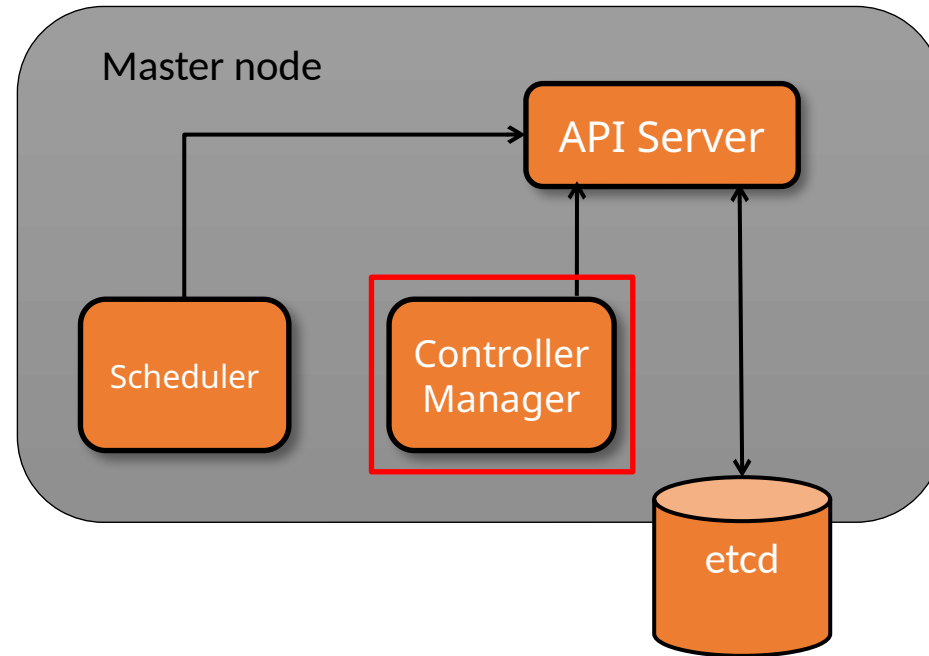
- Scheduler (kube-scheduler): selects nodes for newly created pods to run on





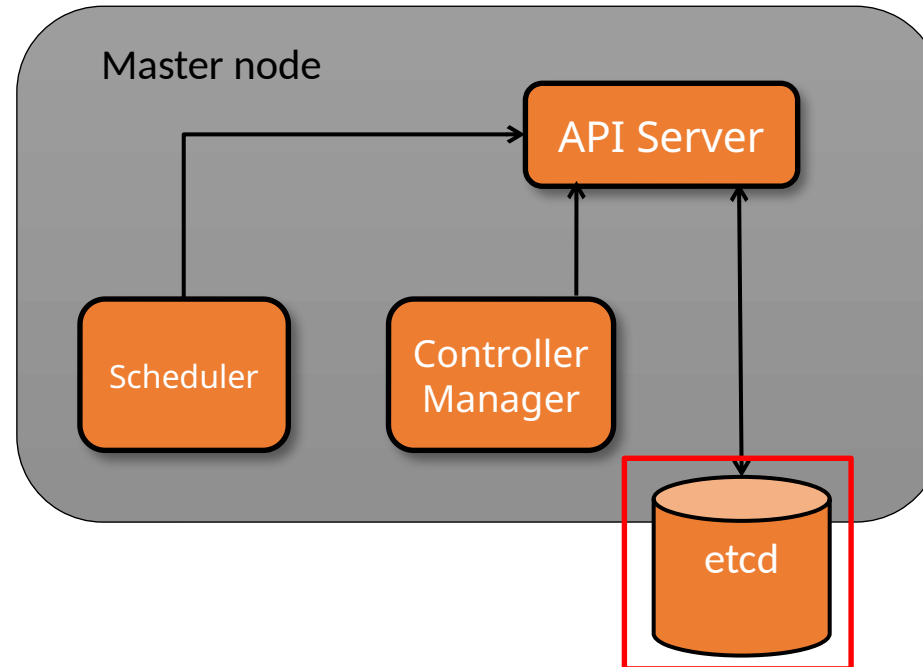
# Master Node Components

- Controller manager (kube-controller-manager): runs background controller processes for the system to enforce declared object states, e.g. Node Controller, Replication Controller,




# Master Node Components

- Persistent data store (etcd): all K8s system data is stored in a distributed, reliable key-value store. etcd may run on separate nodes from the master




# Master Node Deployment Options

- Simple cluster has a single master node
  - At small scale, master may also be a worker node
- Cluster of master node replicas behind a loadbalancer
  - Kube-apiserver and etcd scale out horizontally
  - Kube-scheduler and kube-controller-manager use master election to run single instances at a time



# Worker Node Components


- kubelet: local K8s agent that is responsible for operations on the node, including
  - Watching for pod assignments
  - Mounting pod required volumes
  - Running a pod's containers
  - Executing container liveness probes
  - Reporting pod status to system
  - Reporting node status to system





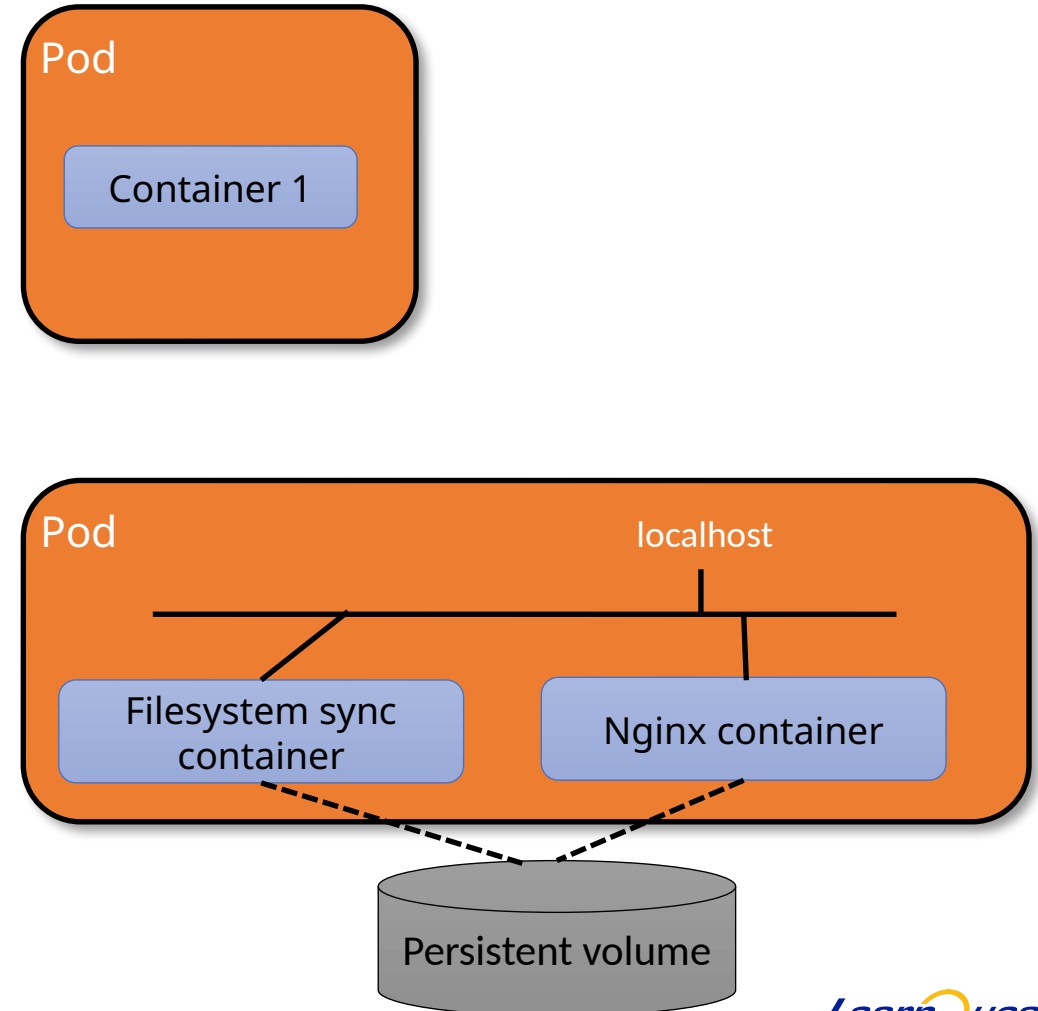
# Worker Node Components

- Service proxy (kube-proxy): enables K8s service abstractions by maintaining host network rules and forwarding connections
- Docker: runs the containers




# Kubernetes Pod

- Basic unit of deployment is the pod, a set of co-scheduled containers and shared resources
- Pods can include more than one container, for tightly-coupled application components, e.g.
  - Sidecar containers : nginx + filesystem synchronizer to update www from git
  - Content adapter: transform data to common output standard
- Containers in a pod share network namespaces and mounted volumes



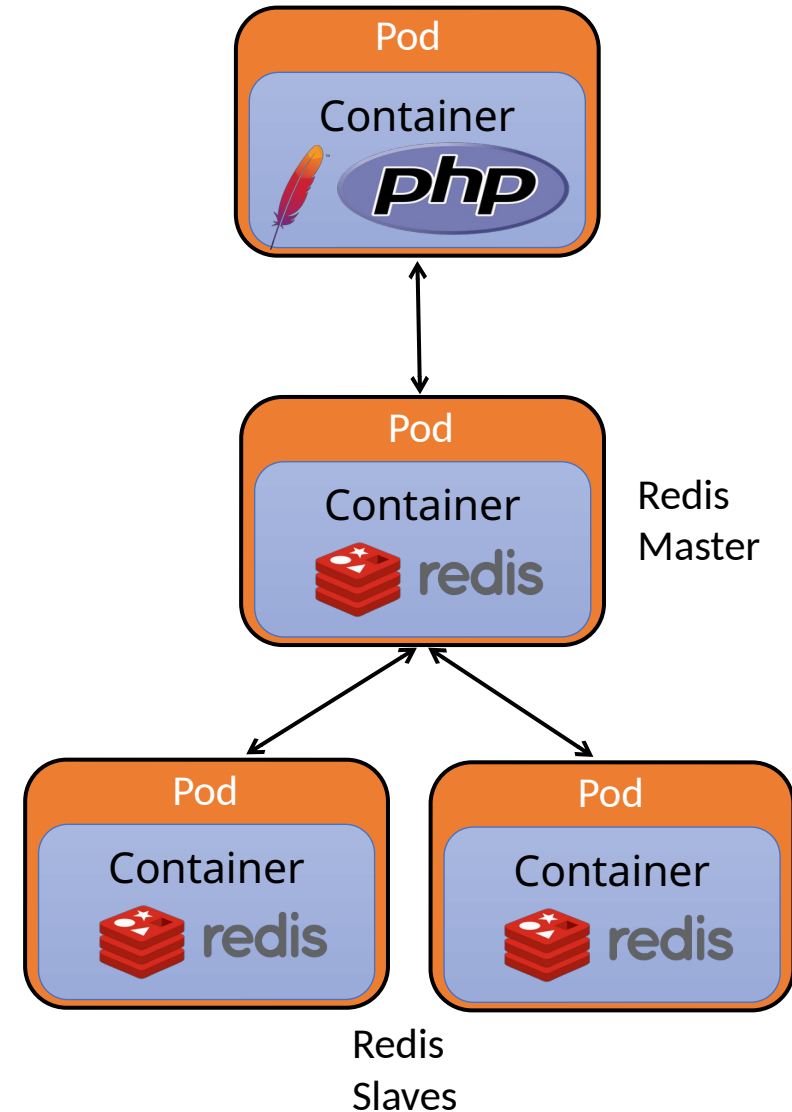
# Pod Deployment

- Possible to use a single pod to run a monolithic application
  - Each application process can be built as a container
  - All containers can access each other's ports on localhost



# Pod Deployment

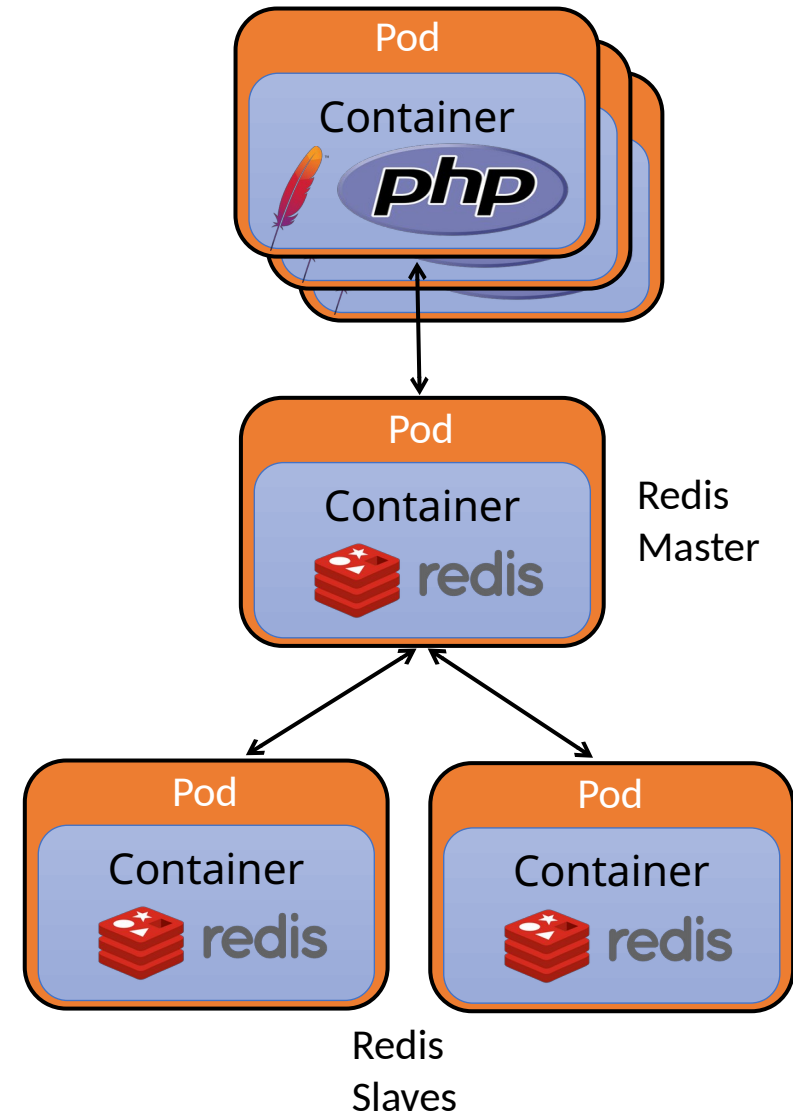
- More advanced features of the K8s system available if application built instead from assemblages of pods, e.g.
  - Web tier: Apache pods
  - Data tier: Redis master/slave pods



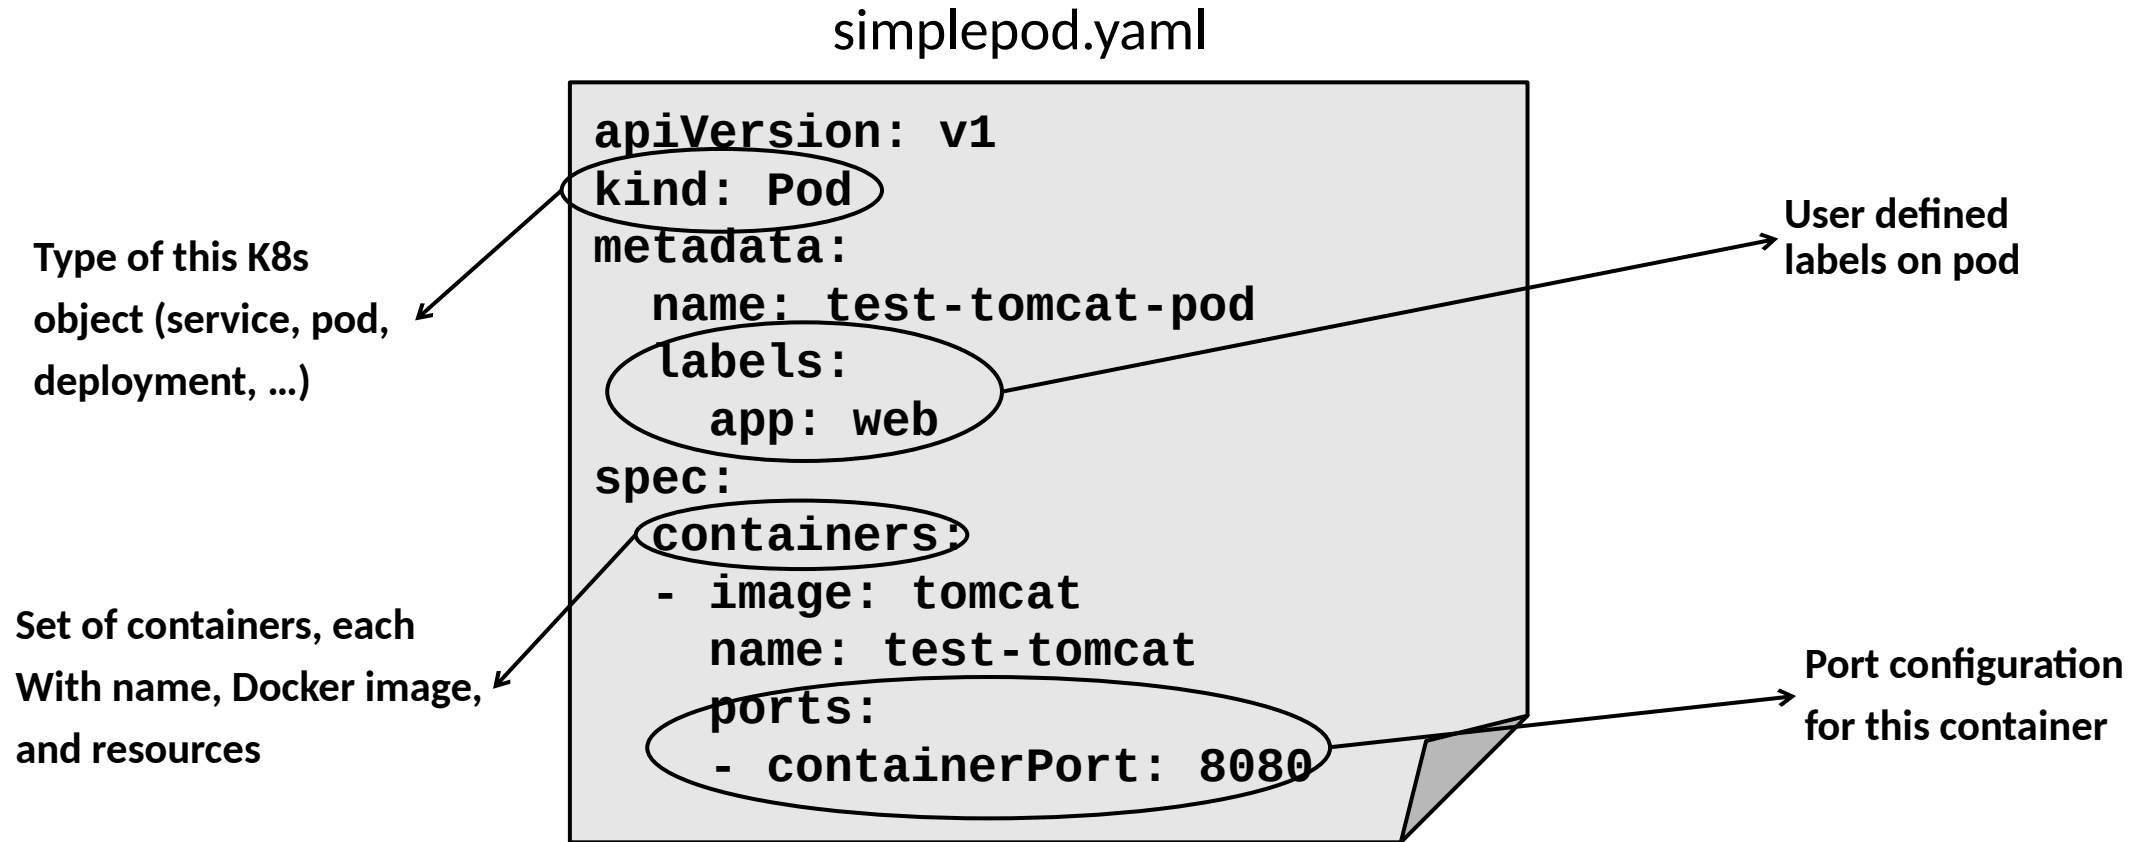


# Pod Deployment

- Pods provide scale and elasticity via replication – not possible in the monolith
- Best practice: assume every pod is mortal



# Defining a Pod via a Manifest File



- Configuration options like creating Docker container directly

# Defining a Pod with Multiple Containers

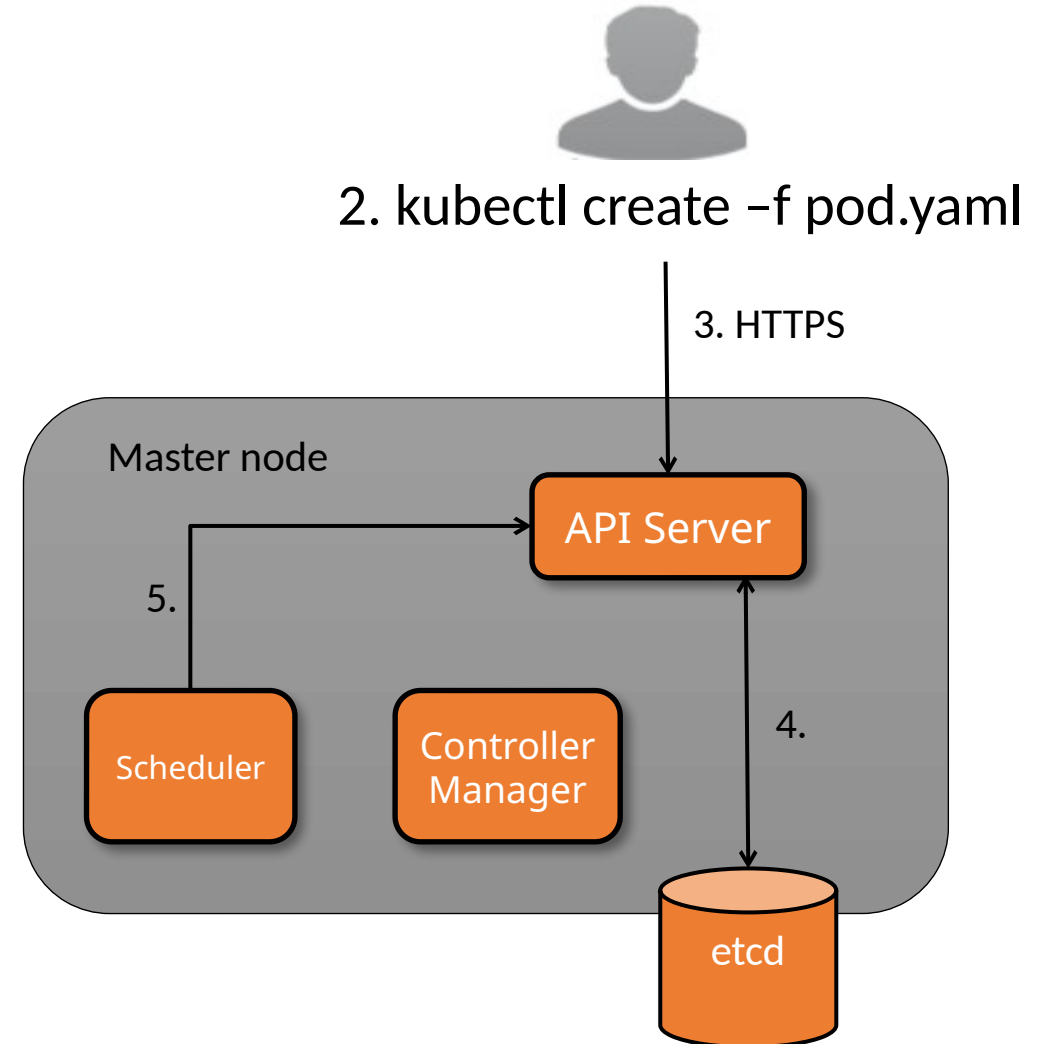
```
apiVersion: v1
kind: Pod
metadata:
  name: test-tomcat-pod
spec:
  containers:
  - image: tomcat
    name: test-tomcat
    ports:
    - containerPort: 8080
  - image: mysql
    name: test-mysql
    ports:
    - containerPort: 3306
```

→ multipod.yaml

- Pod spec can contain multiple containers from different images
- Containers in pod share local network context and cluster IP for pod

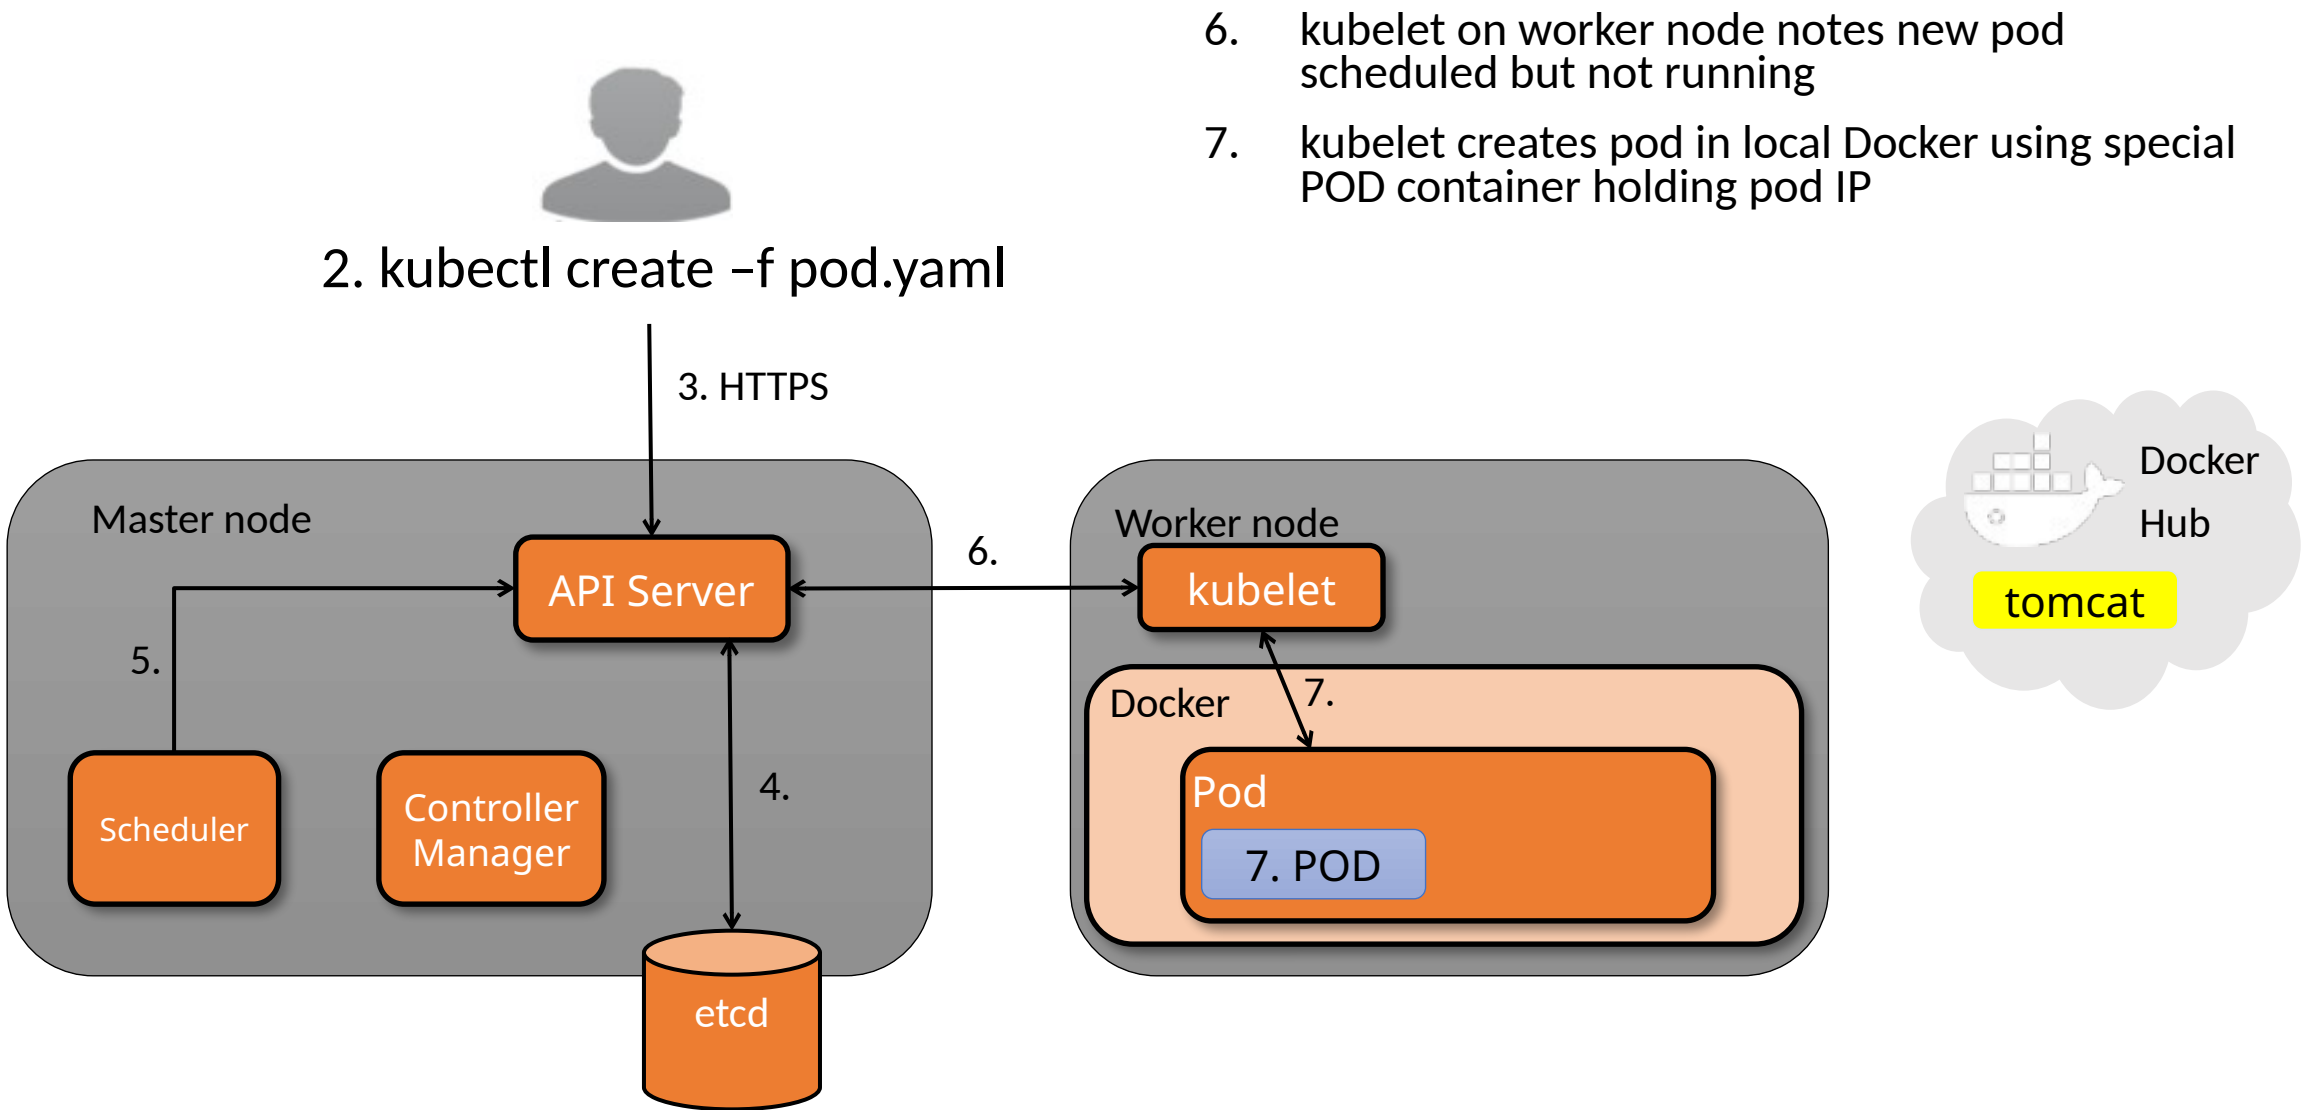
# Pod Creation Process

1. User writes a pod manifest file
2. User requests creation of pod from manifest via CLI
3. CLI tool marshals parameters into K8s RESTful API request (HTTP POST)
4. kube-apiserver creates new pod object record in etcd, with no node assignment
5. kube-scheduler notes new pod via API
6. Selects node for pod to run on
7. Updates pod record via API with node assignment

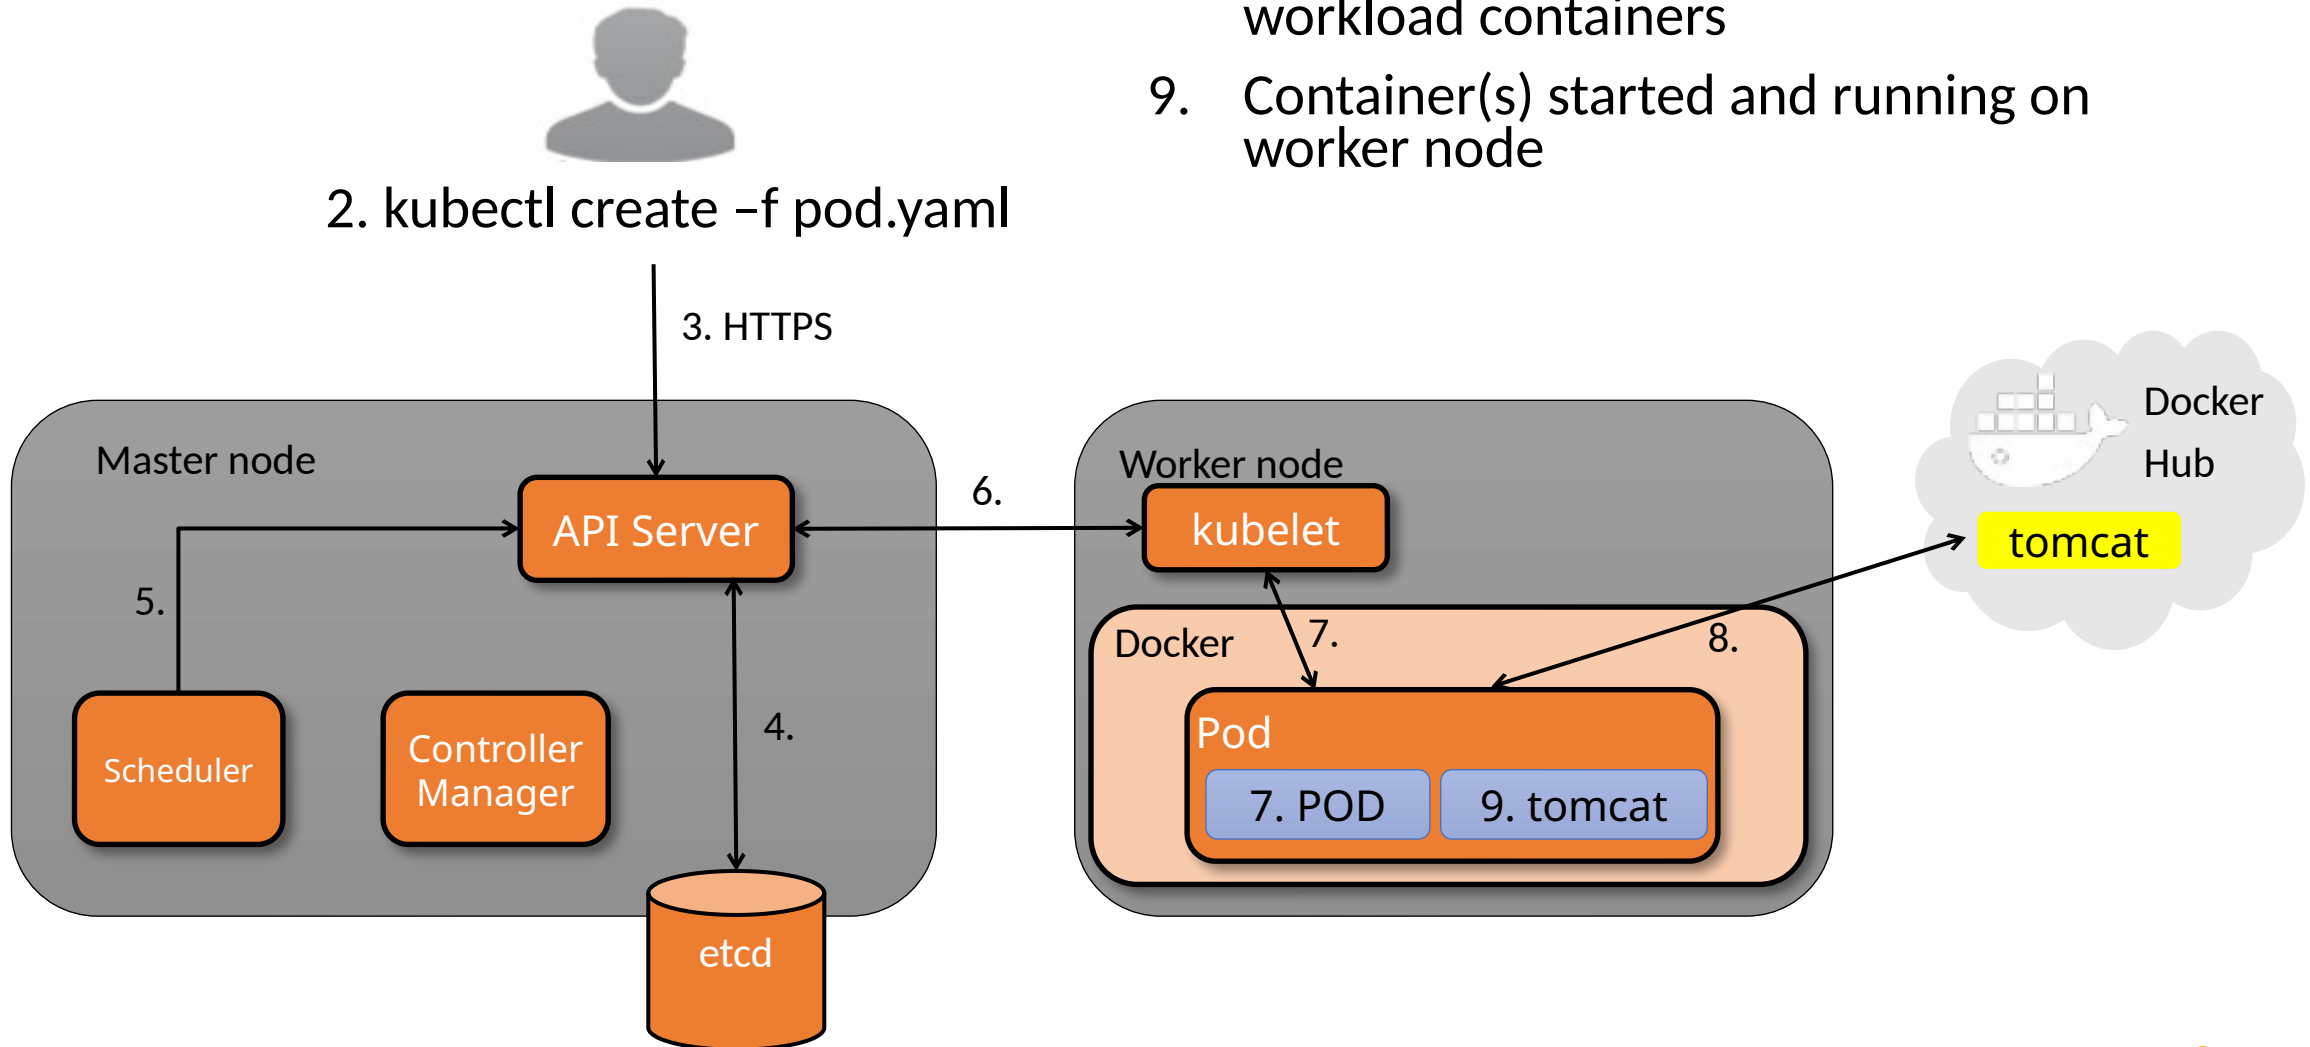




# Pod Creation Process



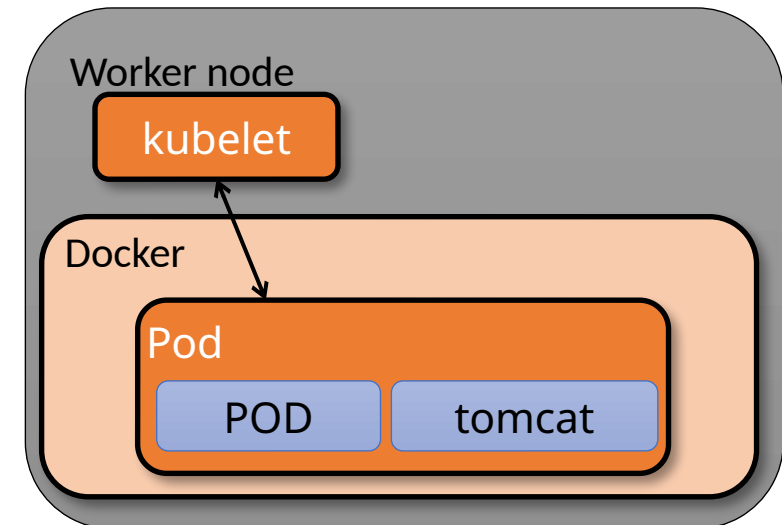
# Pod Creation Process



# Pod Lifecycles


- By default, K8s Pods have an indefinite lifetime, which is not immortality
  - restartPolicy of Always by default
  - restartPolicy of Never or OnFailure also available for terminating jobs
- Node's kubelet will create and keep running containers for pods assigned to node, per the pod specs
- If a Pod container fails to start, or unexpectedly exits, kubelet will restart it
  - Can see container lifecycle events via 'kubectl describe pod <PODNAME>'
- If node is lost, its Pods are also lost – K8s will not rebind Pods to another node

```
apiVersion: v1
kind: Pod
metadata:
  name: test-tomcat-pod
  labels:
    app: web
spec:
  containers:
  - image: tomcat
    name: test-tomcat
    ports:
    - containerPort: 8080
```



# Deleting Pods



- When deleting a Pod, its containers will be removed and pod IP relinquished
- If an application needs to persist data, its pods must be configured to use persistent volumes for storage
- If a node dies, its local pods are also gone
- Best practice: use controller resources instead of managing pods directly
- Best practice: use service resources to build reliable abstraction layers for clients





# Kubernetes Pods

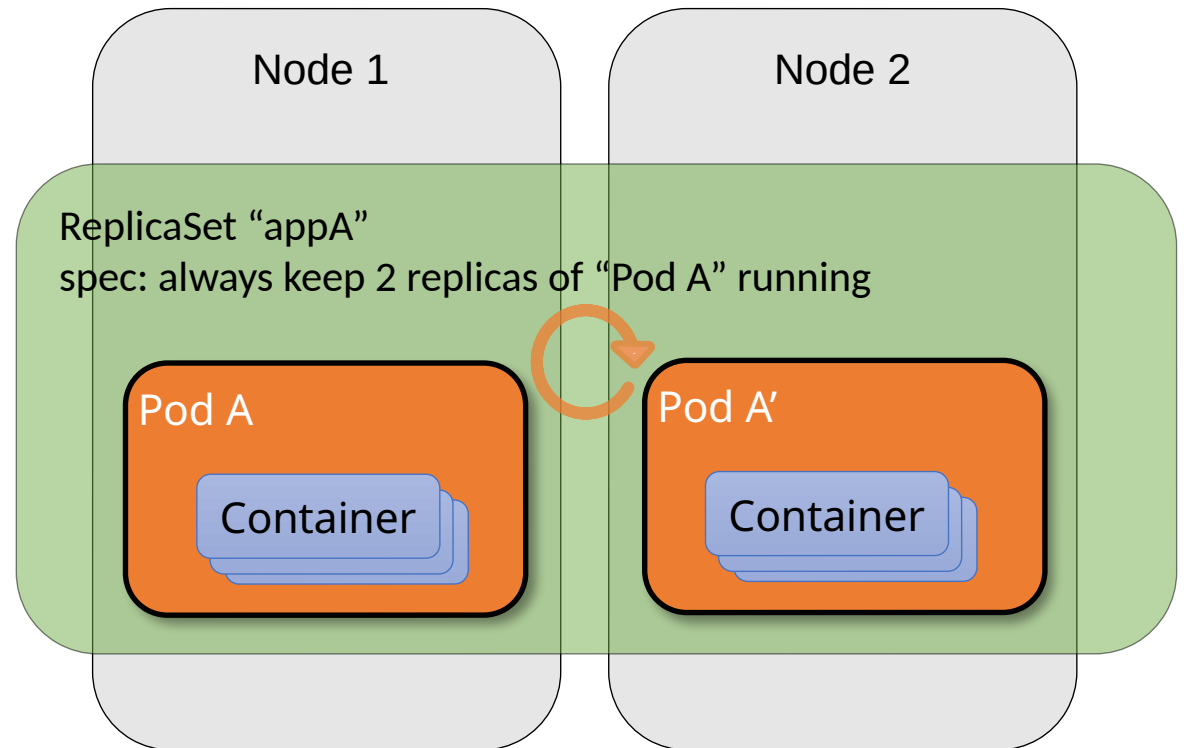
- Pod is an application instance
  - Pods can include more than one container, for tightly-coupled application components
  - Containers in the same Pod share networking and storage resources
- Kubernetes handles efficient placement of Pods across available Nodes



# Application Patterns

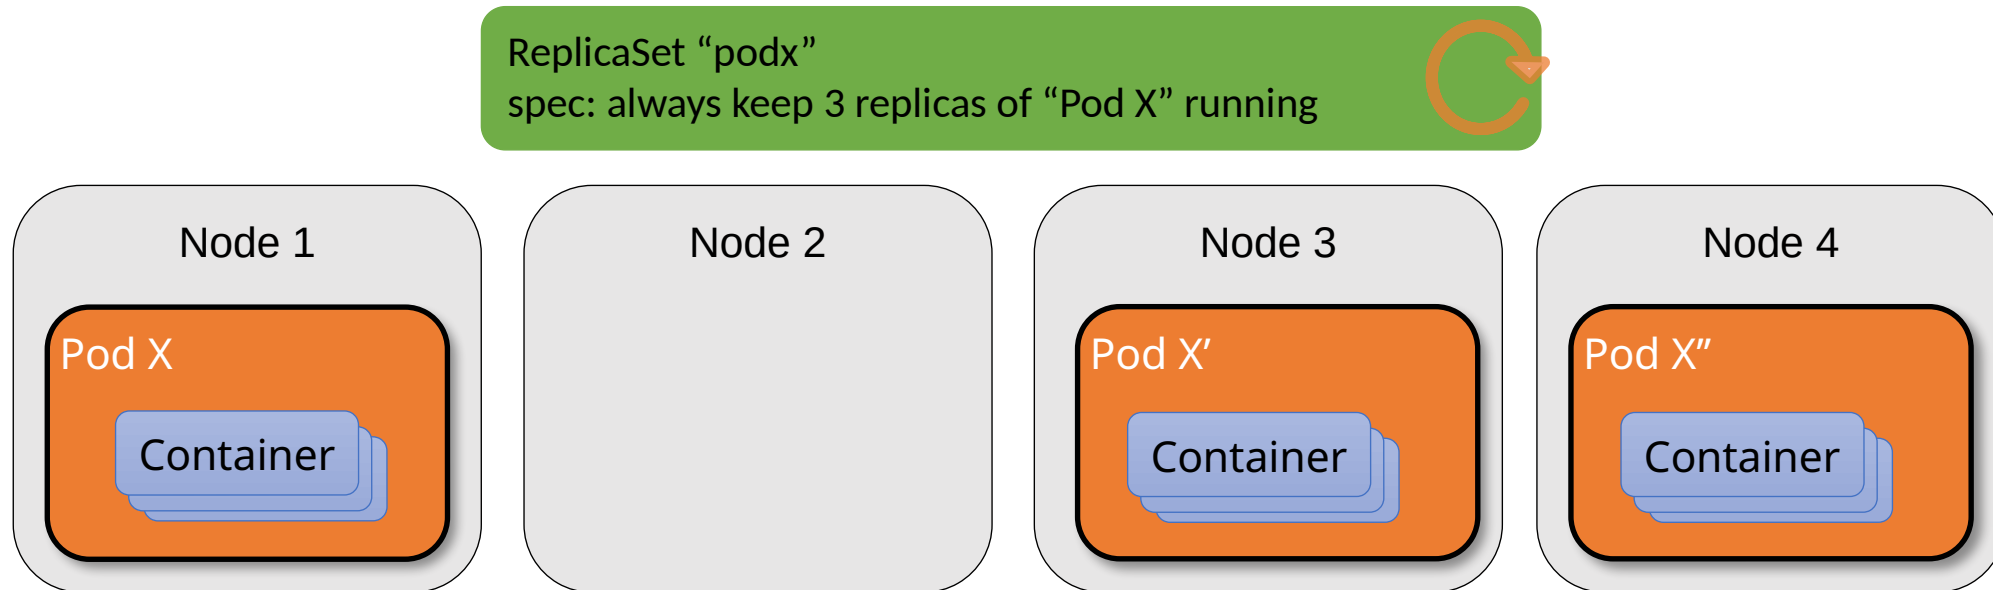
- K8s controller creates and manages Pods according to different application
  - ReplicaSets manage sets of replicas of stateless workloads to ensure availability
  - StatefulSets manage stateful workloads on stable storage to ensure consistency
  - DaemonSets manage workloads that must run on every node, or set of nodes
- Jobs manage parallel batch processing workloads

## Controller example: ReplicaSet



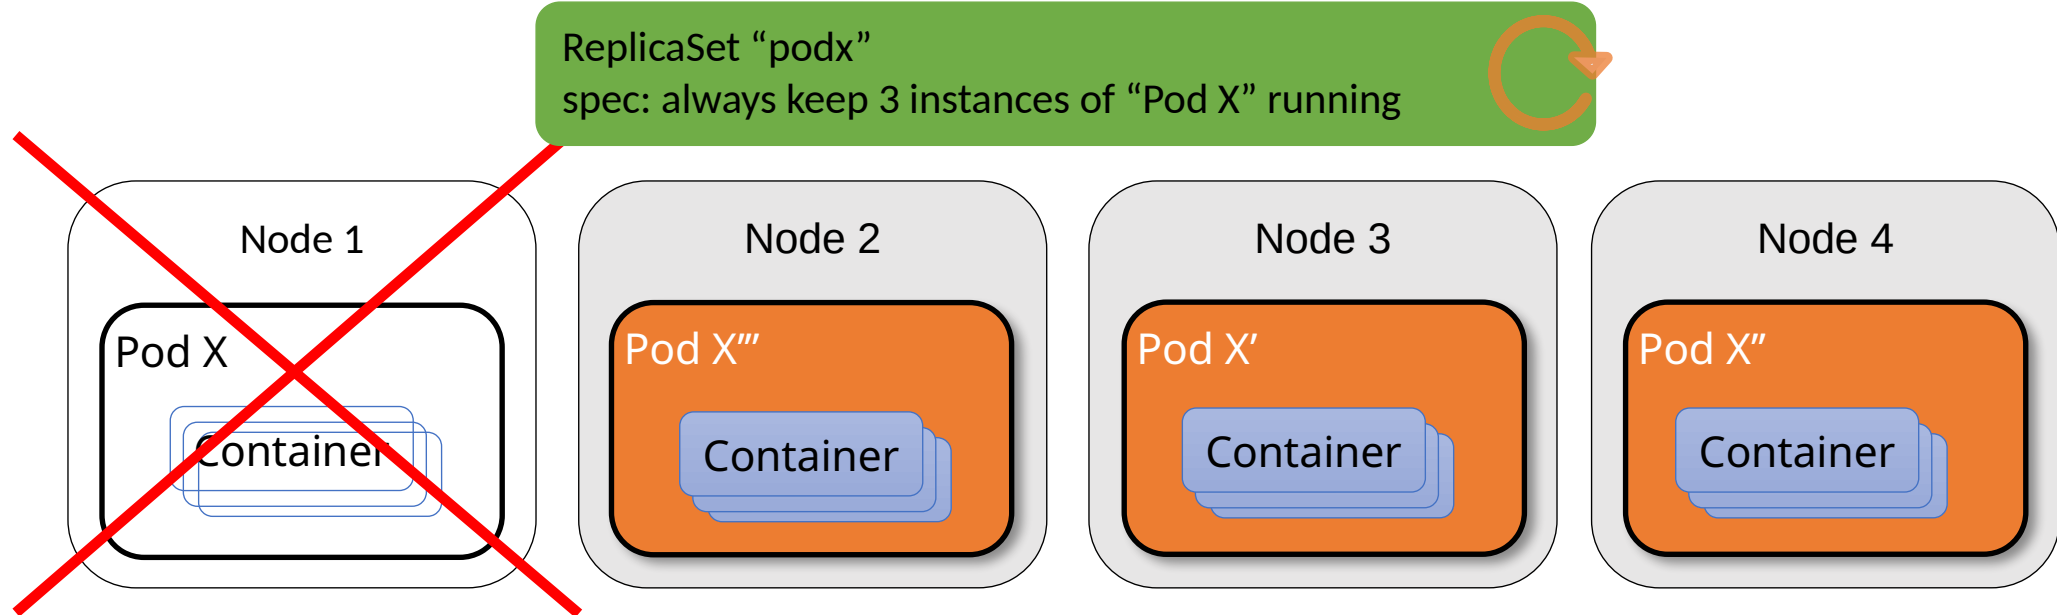
# ReplicaSets

- ReplicaSet configuration specifies how many instances of given Pod exist
- ReplicaSet used for web applications, mobile back-ends, API's



# Replication Ensures Application Availability

- When a Node fails, its Pods are lost
- K8s system manages the state of the ReplicaSet back to the declared configuration
- Changing the configuration will result in management to new state, e.g. scale out



# Objects and Resources

- Objects are the persistent entities that users manage via the Kubernetes API
- Objects track what's running and where, available system resources, and behavioral policies, e.g

## Workloads

- Pod (run)
- Service (expose)

## Configuration

- Secret
- ConfigMap

## Controllers

- ReplicaSet
- Deployment
- StatefulSet
- DaemonSet
- Jobs / Cron Jobs

## Workload Persistent Storage

- PersistentVolume
- PersistentVolumeClaim


## Cluster Resources

- Node
- Namespace
- Cluster

## Network Resources

- Ingress
- NetworkPolicy



# The Essential Architecture Review





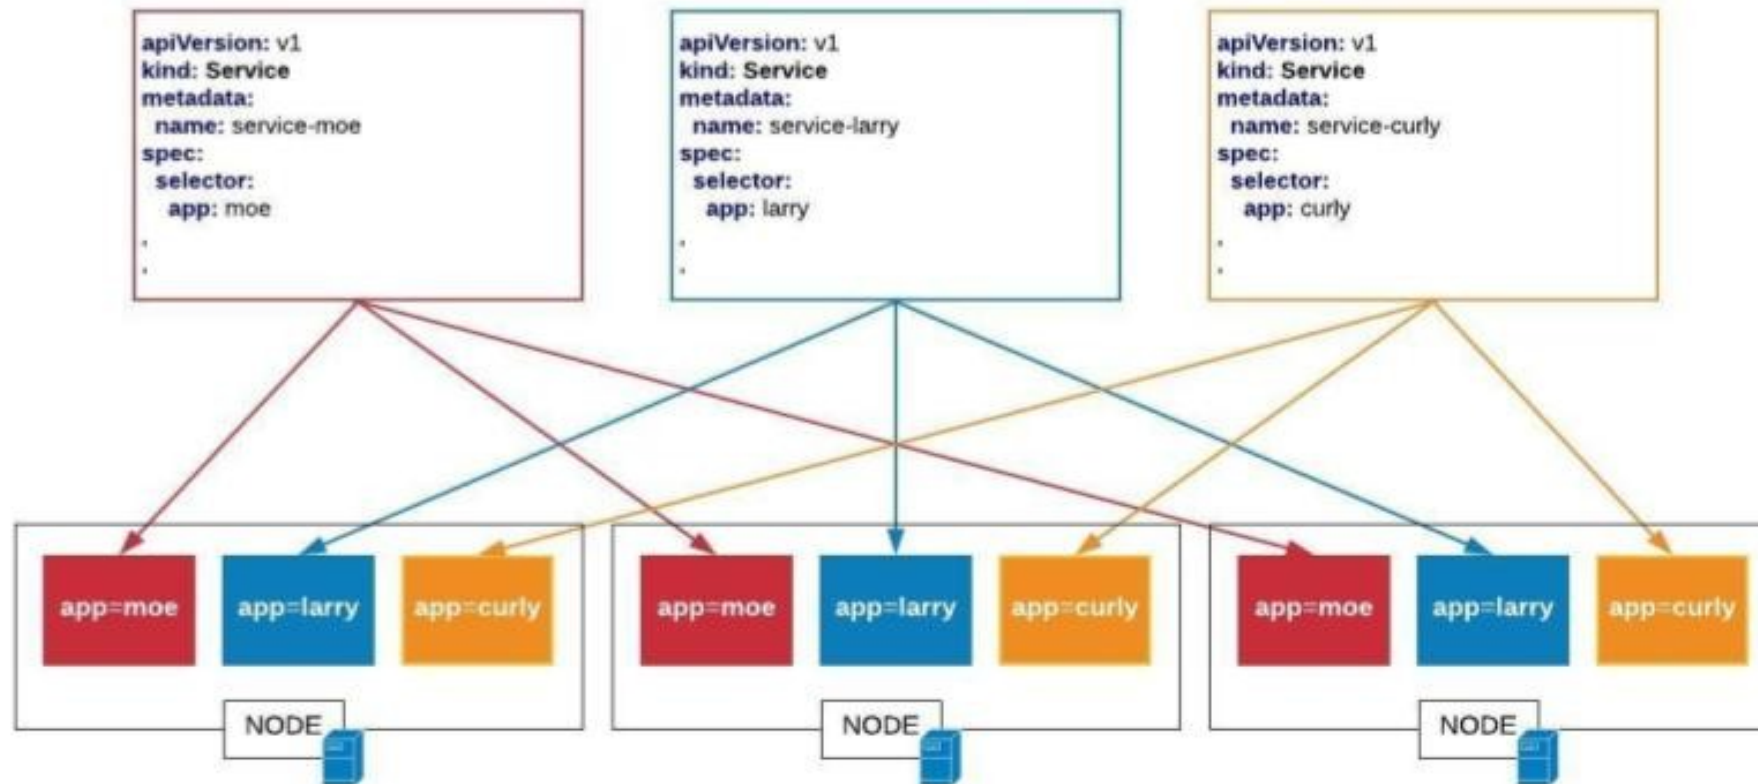
# The Essential Architecture Review

## Pod




# Services and Service Binding

## Service




# Services and Service Binding

## Service - ClusterIP




# Services and Service Binding

## Service - NodePort

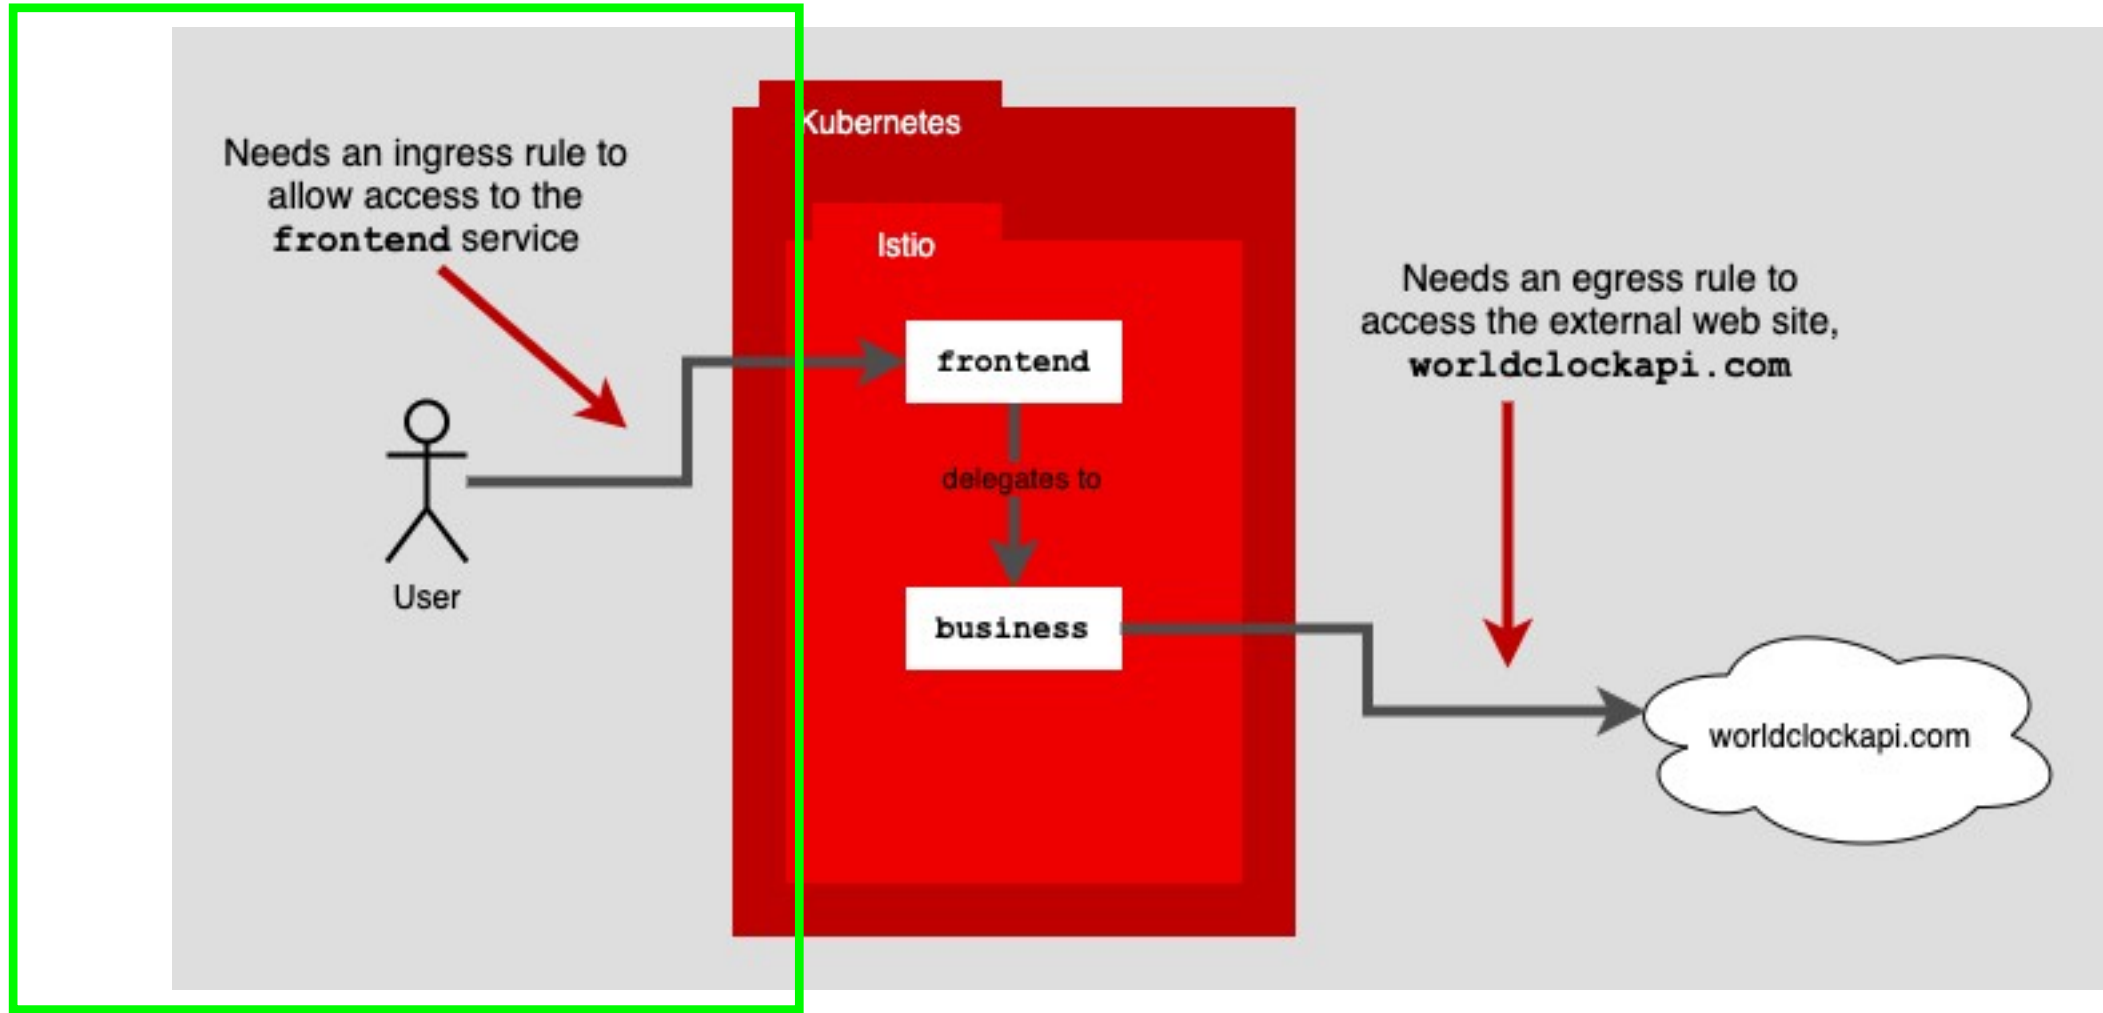


# Services and Service Binding

## Service - LoadBalancer

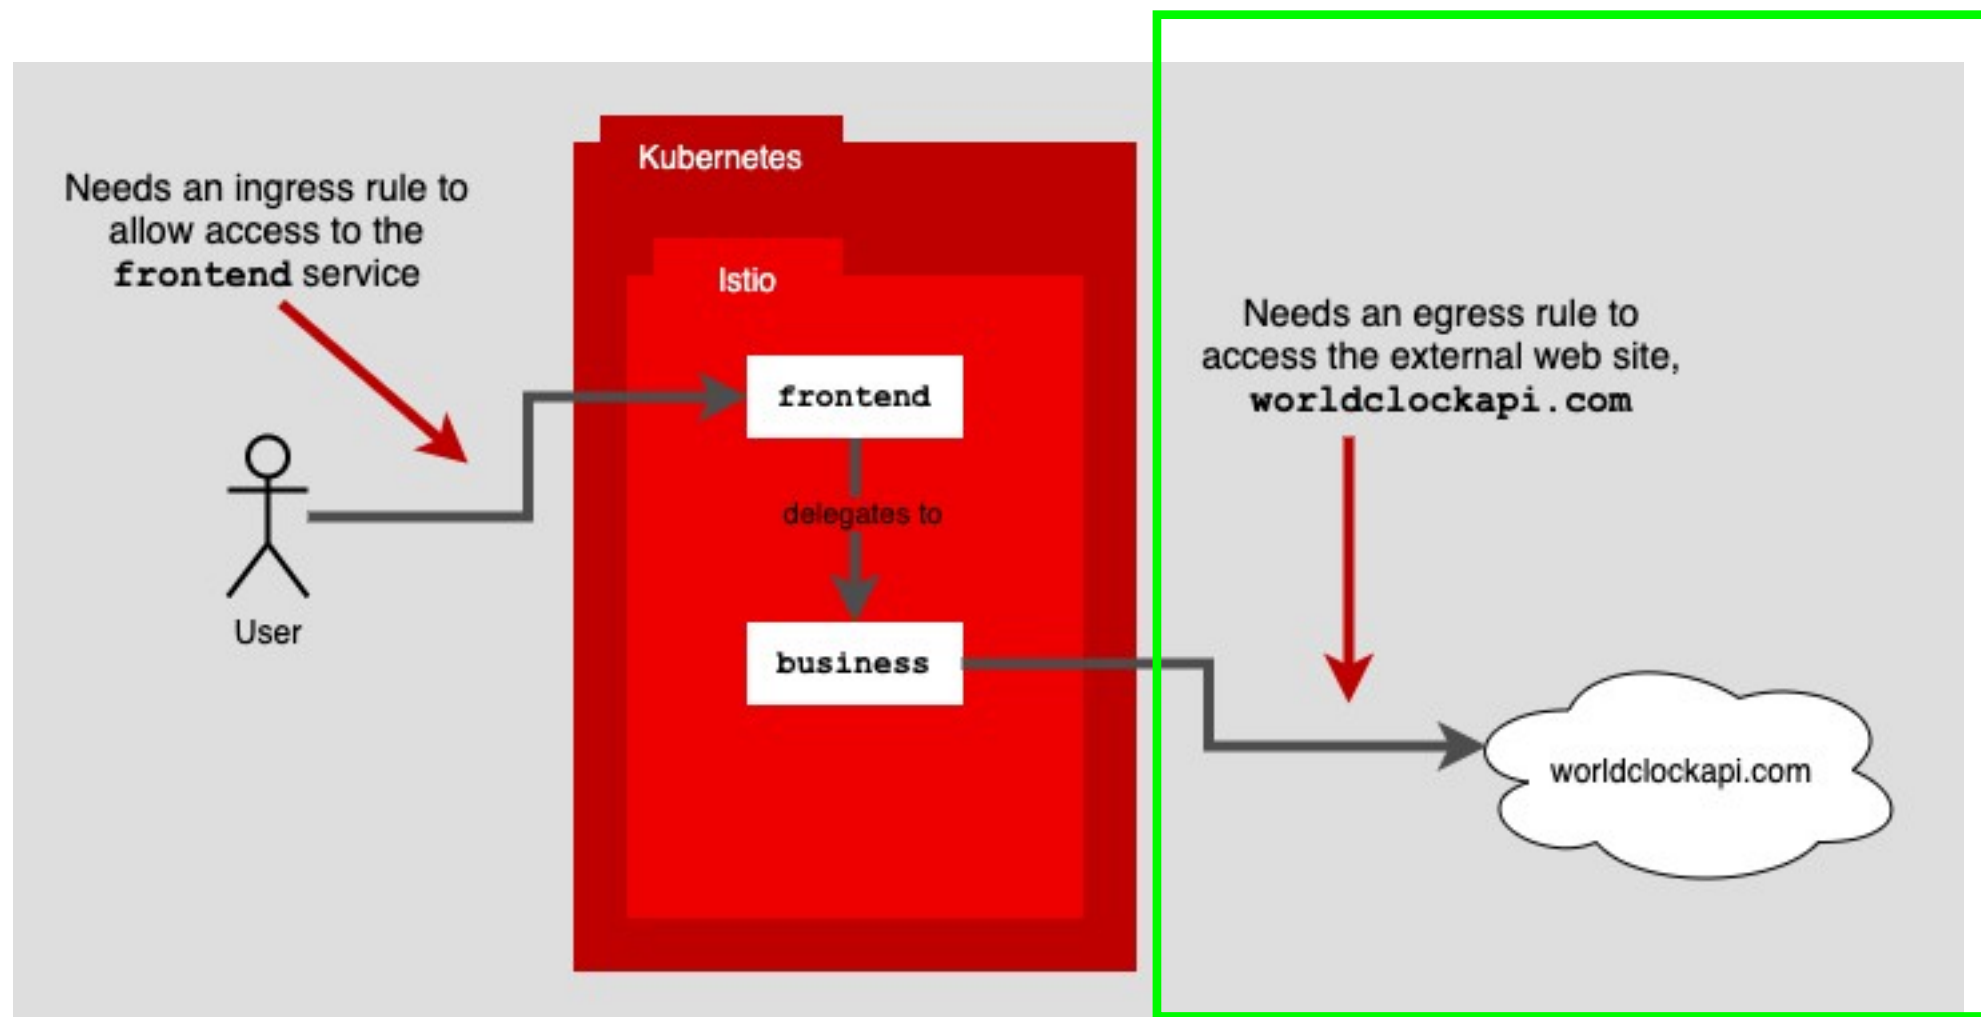


# Kubernetes: Ingress

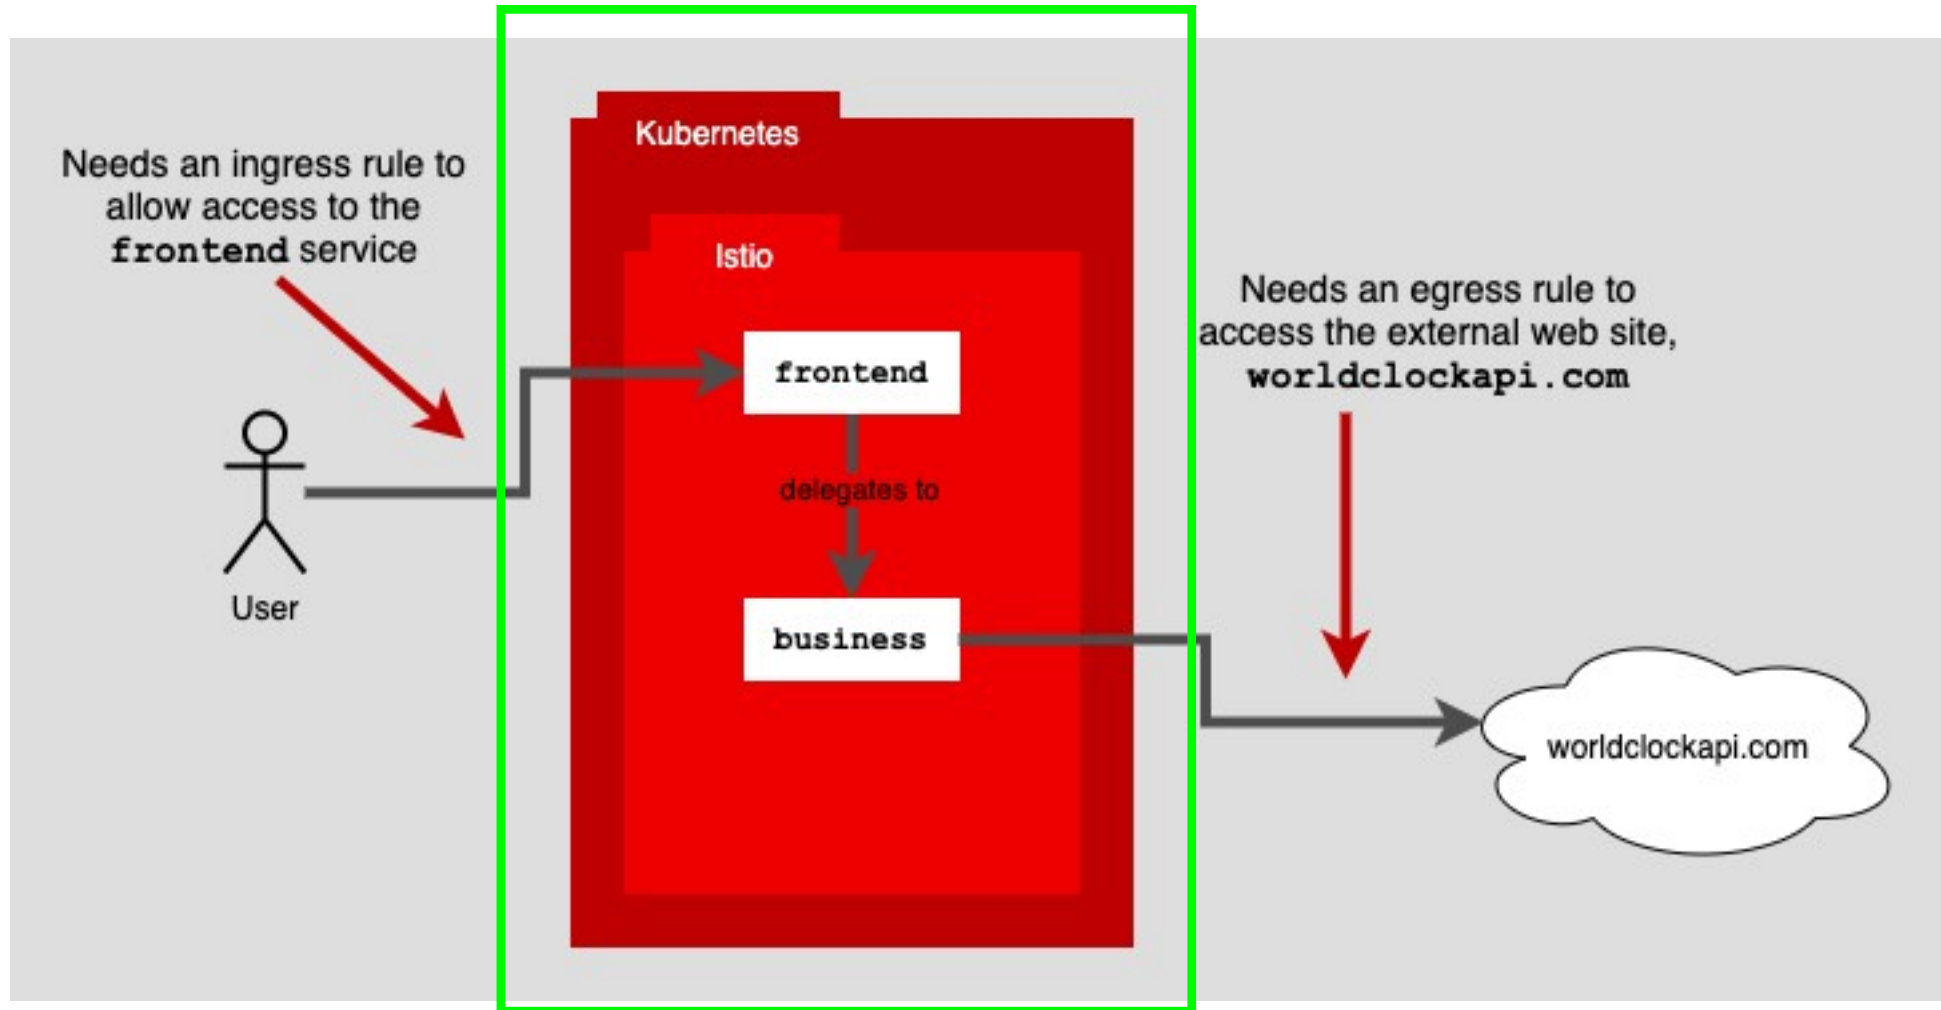




# Kubernetes: Egress




# Service Mesh



# Working with a Service Mesh

- What does a Service Mesh do?
  - Enforces routing rules and restrictions
  - Binds service to implementation instance
  - Determines optimal instance
  - Keep track of performance aspects, e.g., latency
  - Coordinates retries and failures
  - Ejects failing instances from the load balancing pool
  - Provides monitoring, e.g. tracing and performance metrics

# Sampling of Products



Service Fabric Mesh!



A detailed Renaissance-style painting of Plato's Academy. The scene is set in a grand hall with classical columns and a landscape view through the arches. Numerous figures, representing various Greek philosophers, are engaged in discussion and study. Plato, an older man with a white beard wearing a red robe, stands in the center, pointing his right index finger towards the sky. Aristotle, a younger man with a dark beard in a brown robe, stands next to him, gesturing with his palm facing down towards the earth. Other figures are seated or standing, some holding books or scrolls. The overall atmosphere is one of intellectual pursuit and philosophical discourse.

**Questions?**



# End Module

